

Red Hat OpenShift Container Platform


- Public/private DC.
- Bare metal and multiple cloud and virtualization providers.
- Full control by customer.

Red Hat OpenShift Dedicated



- Managed cluster in public cloud.
- RH manages the cluster.
- Customer manages updates and add-on services.

Red Hat OpenShift Online


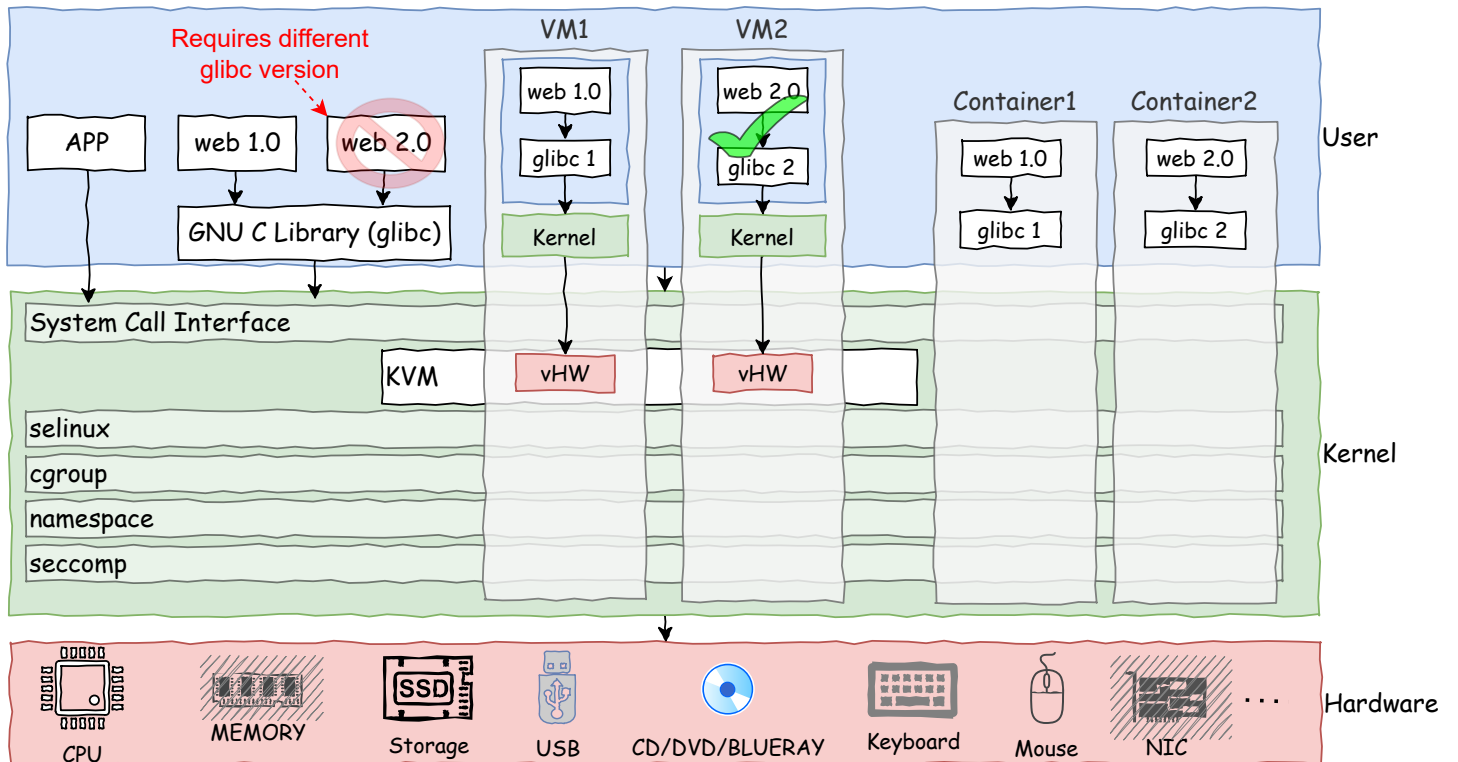
- Public hosted cluster.
- Shared resources by multiple customers.
- RH manages cluster life cycle.



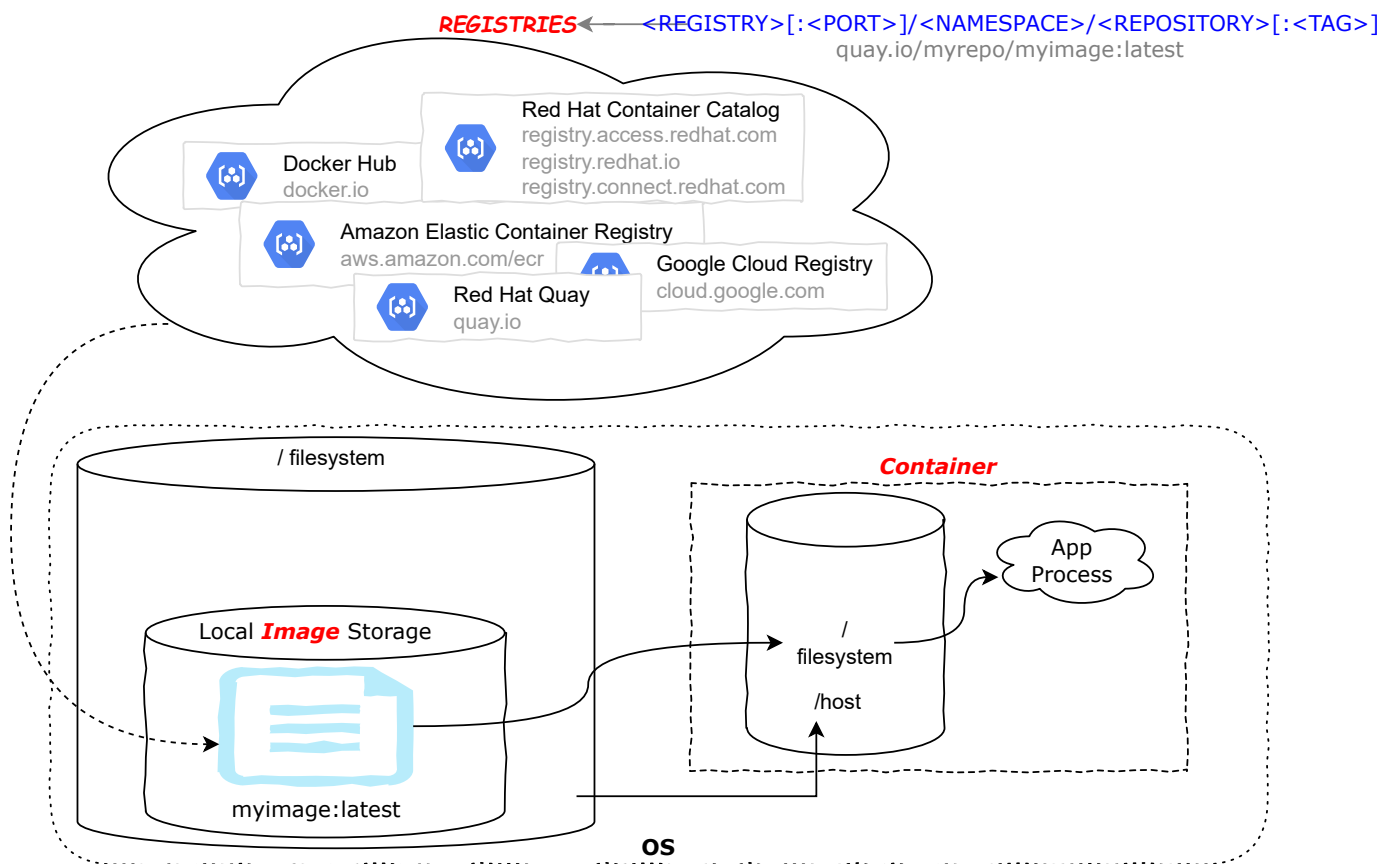
Kubernetes Declarative Architecture



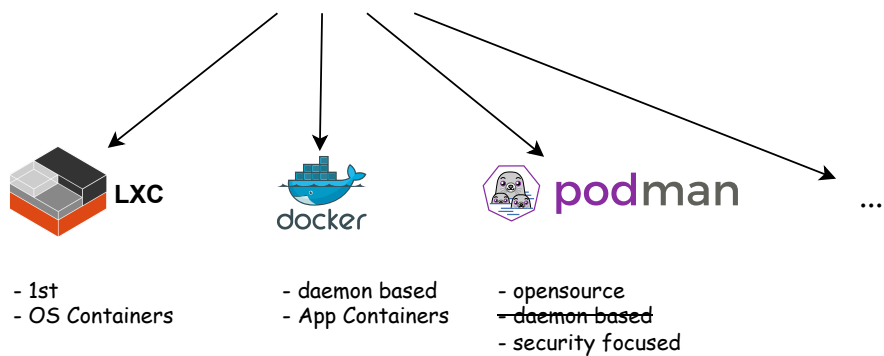
VM vs Container



Container Architecture



Container Utilities



OS Container Vs Application Containers

Podman

Image and Registry Operations

`podman login [-u USER] [-p PASS] [REGISTRY]` Only if required. Accessing private repo or updating image.

`podman logout {-a | REGISTRY}` Logout of registry (-a for all).

`podman images [-q]` List local images (-q only show id).

`podman rmi IMAGE...` Removes local image(s). Use -af with caution.

`podman search KEYWORD` Search registry for an image.

`podman pull SOURCE` Pull image from a registry.

Where,

SOURCE [*REGISTRY*[:*PORT*]/*NAMESPACE*/]*IMAGE*[:*TAG*]
dir:*PATH*
docker-archive:*PATH*
oci-archive:*PATH*

`podman tag IMAGE[:TAG] TARGET_NAME[:TAG]` Add an additional name to a local image

`podman push IMAGE` Upload an image to the registry

Container Operations

`podman run [--name NAME] [-p PORT_INFO] [-v VOL_INFO] [-d] [-it] IMAGE [CMD_INFO]`

Where,

--name *NAME* Container name. Autogenerated if not provided.

-p *PORT_INFO* [*LOCAL_IP* :] *LOCAL_PORT* [[: *CONT_IP*] : *CONT_PORT*]
Mapping between local IP:PORT to container IP:PORT

-v *VOL_INFO* *LOCAL_DIR* : *CONT_DIR*
Mapping between local dir to container dir.

-d Run in detached mode (background).

-it -i keep stdin open, -t allocate a pseudo-tty.

IMAGE Image used to create the container.

CMD_INFO *CMD* [*ARG...*]
Command to run in container.

`podman exec [-it] CONTAINER CMD_INFO` Execute command inside running container.

`podman ps [-a] [-q]` List containers (-a for all, -q only show container id).

`podman rm CONTAINER...` Remove one or more stopped containers.
(-f includes running and paused containers).

`podman start|stop|restart CONTAINER...` Start, stop or restart one or more containers.


`podman kill [-s SIGNAL] CONTAINER...` Send signal to one or more containers.

For more info:

`podman --help` OR `man podman`

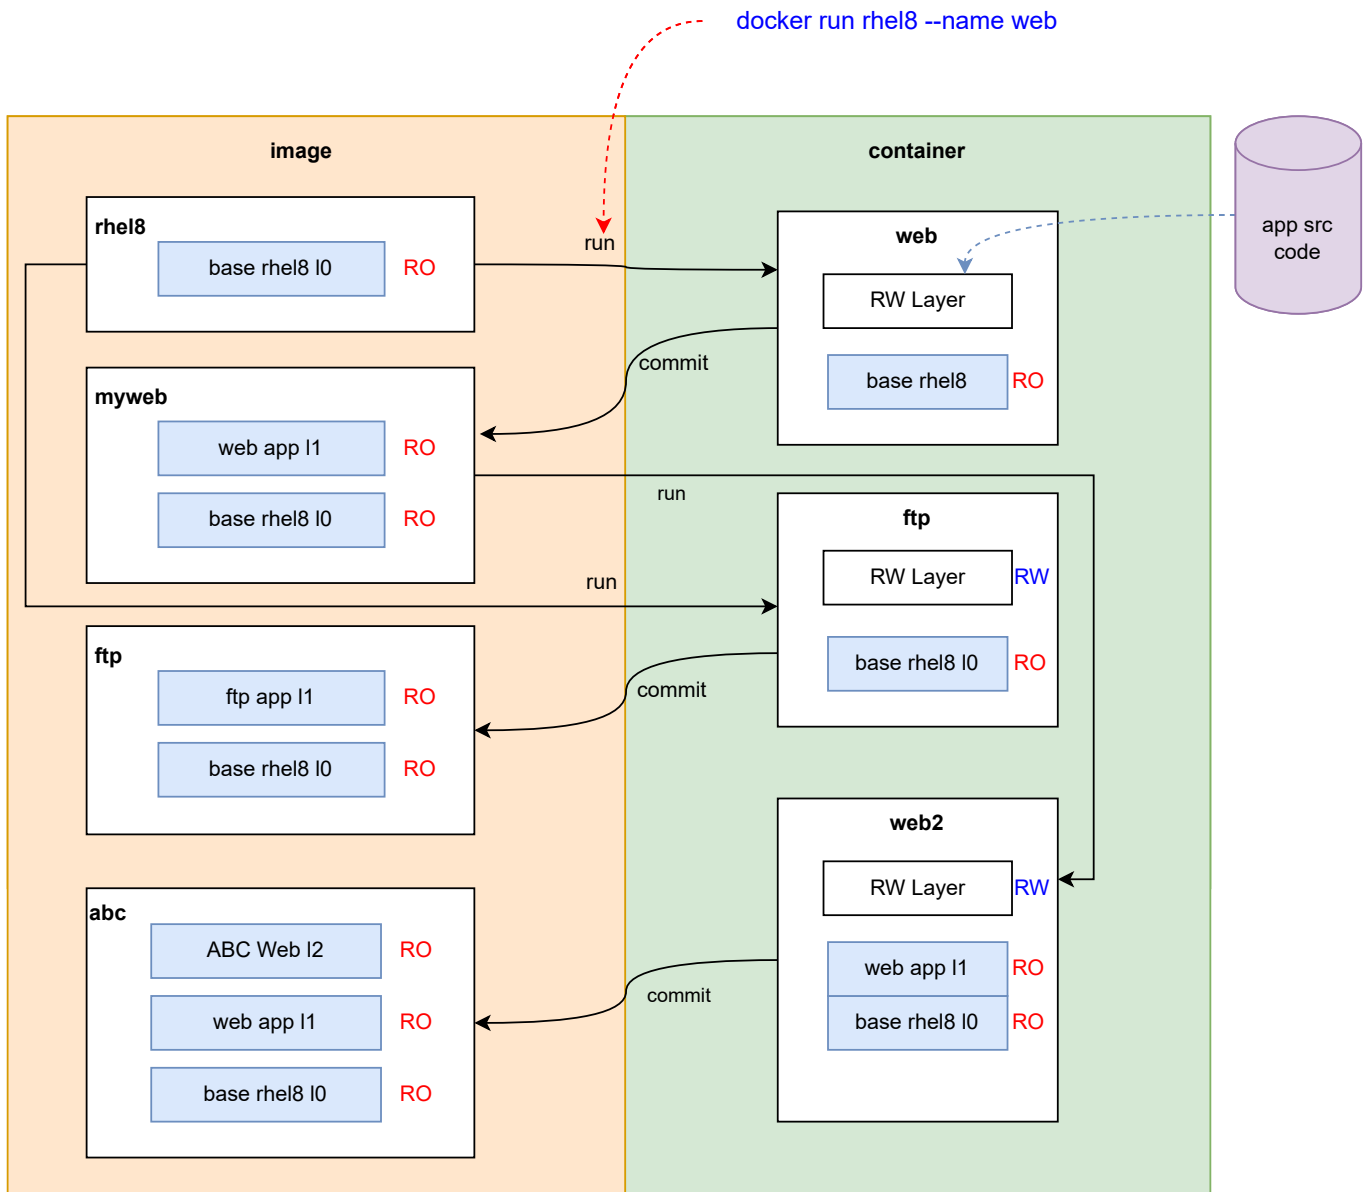
Each sub command has it's own man page. i.e `man podman-run`, `man podman-images`, etc.

Basic Container Design

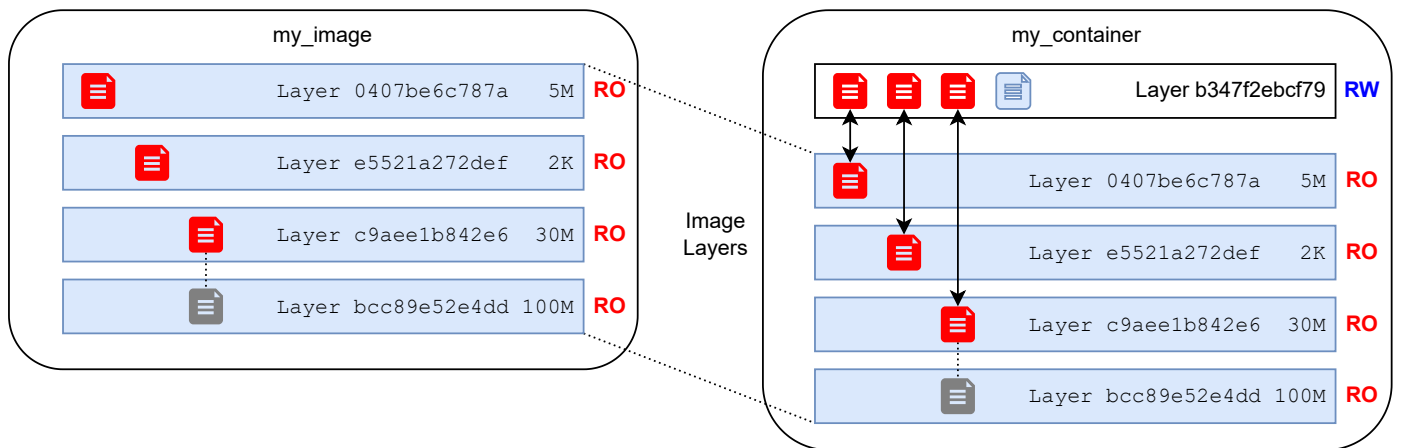


Creating Image

1. Manual
2. Dockerfile/Containerfile
3. Source-To-Image(s2i/STI)
 - a) get runtime image and create container
 - b) clone source code into container
 - c) compile source code
 - d) deploy/publish compiled app
 - e) cleanup
 - f) save container as image




UnionFS - A Stackable Unification File System



BASE IMAGE TYPES


MINIMAL



Designed for apps that contain their own dependencies (Python, Node.js, .NET, etc.)

- Minimized pre-installed content set
- no suid binaries
- minimal pkg mgr (install, update & remove)


PLATFORM



For any apps that runs on RHEL

- Unified, OpenSSL crypto stack
- Full YUM stack
- Includes useful basic OS tools (tar, gzip, vi, etc)

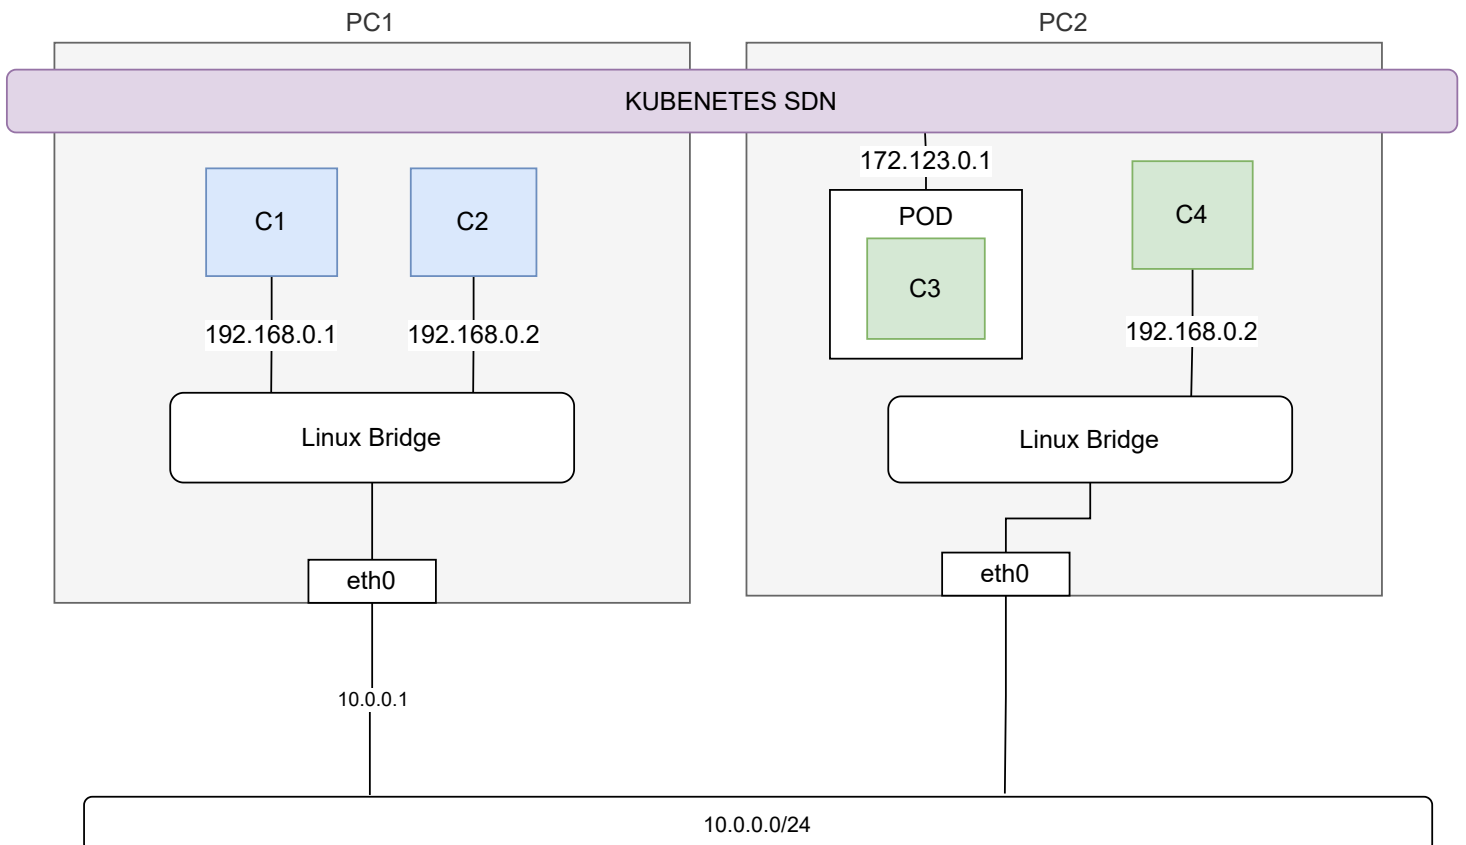
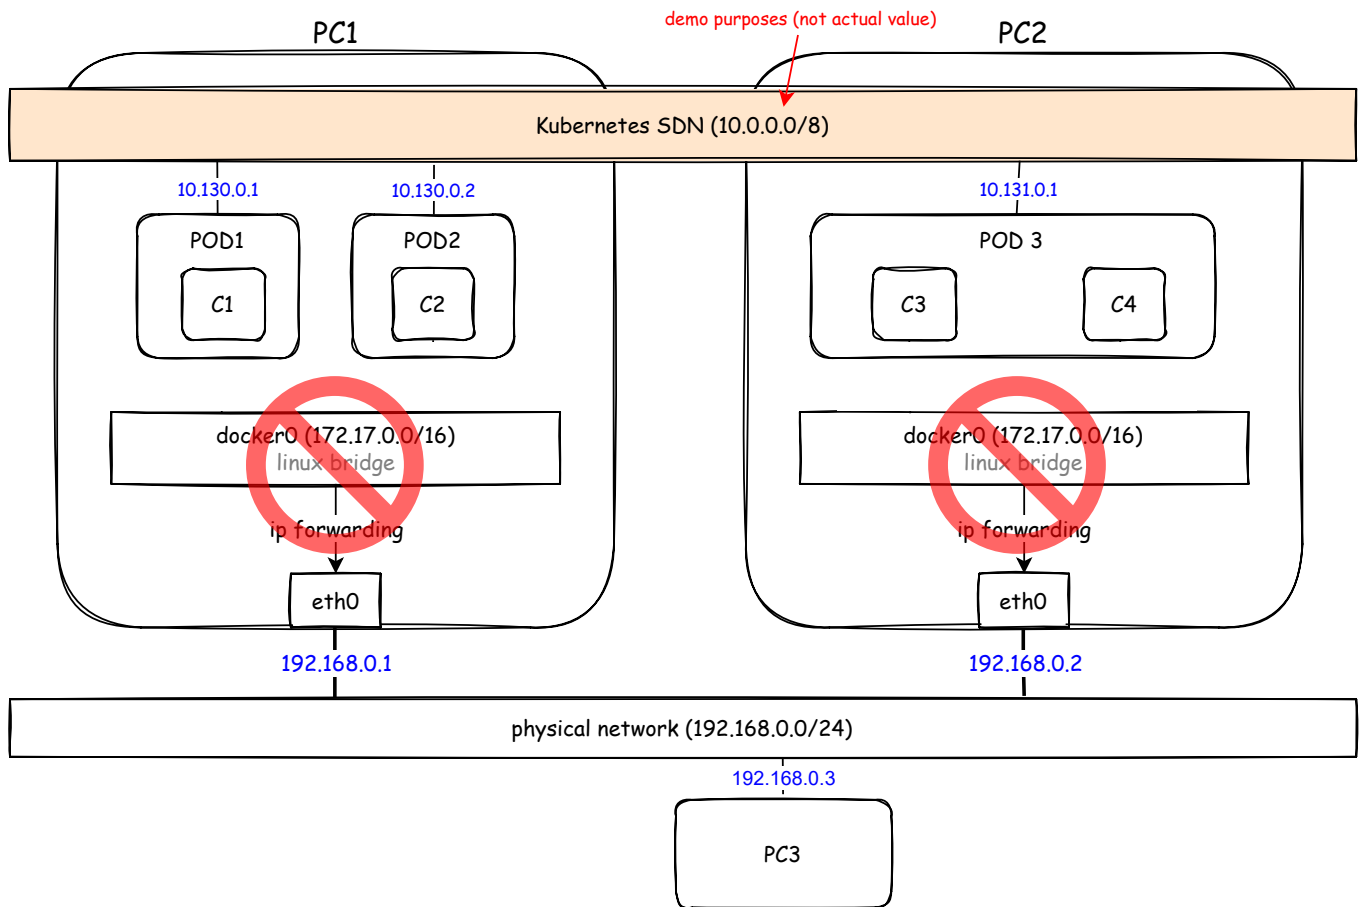
MULTI-SERVICE



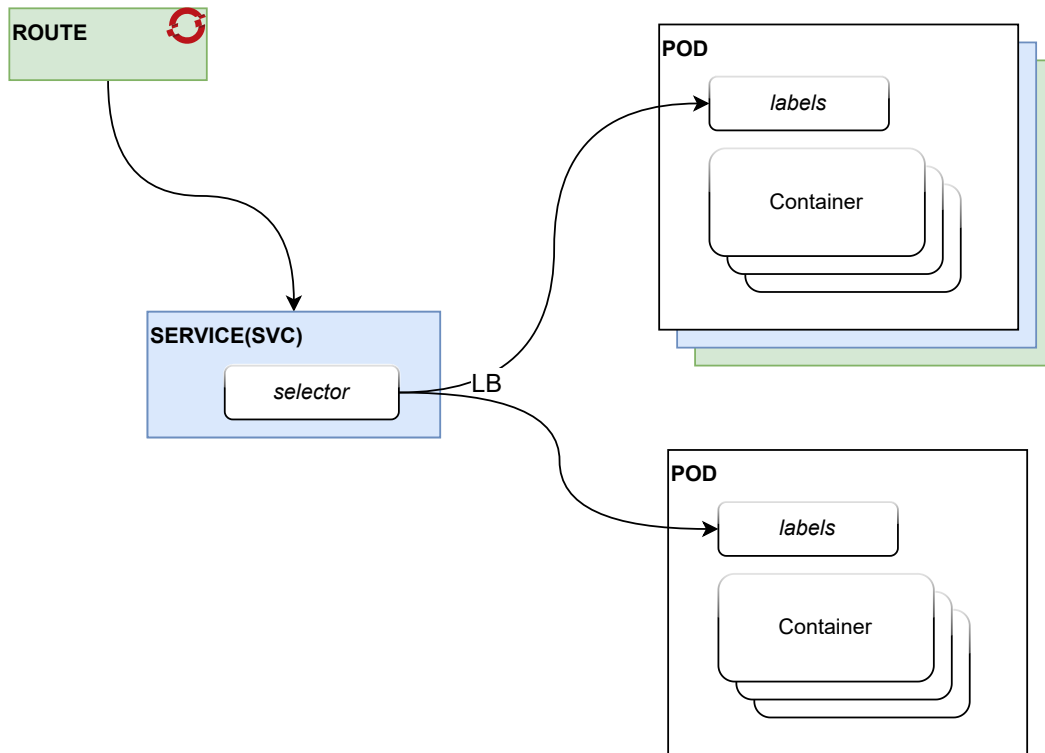
Eases running multi-service in single container

- configured to run systemd on start
- allows you to enable th services at build time

Basic Network - Container vs Kubernetes



Route, Service and Pod Relationship



POD

A pod contains one or more containers.

SERVICE

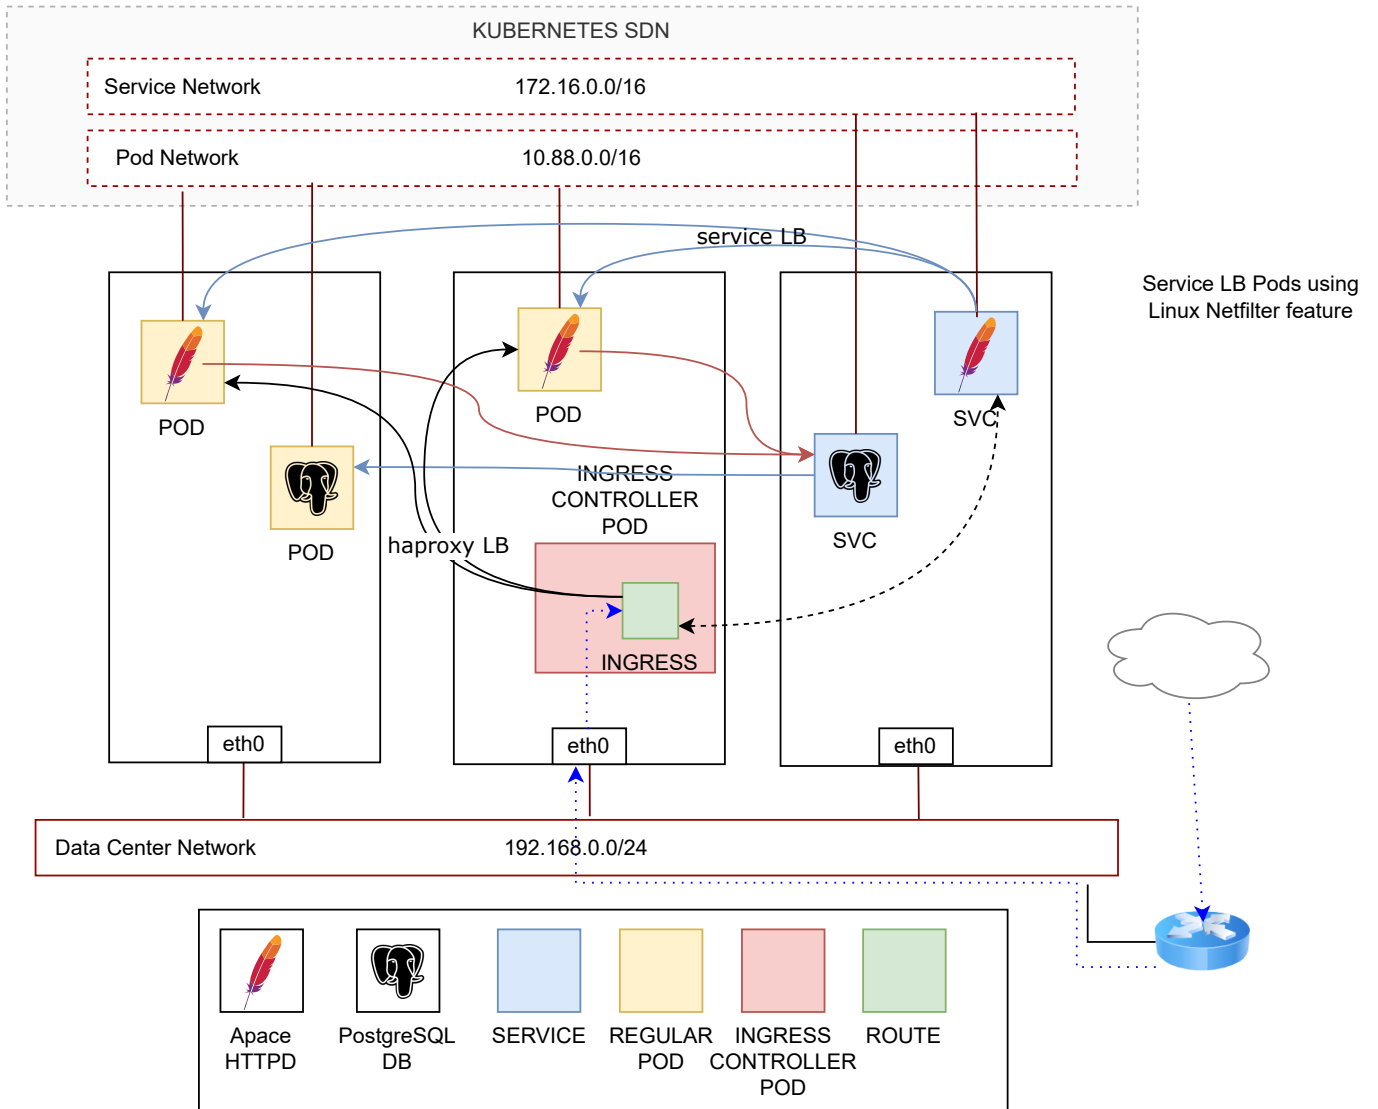
A service references the pod(s) by using the label selector.
The service load balances the connections between all the pods.

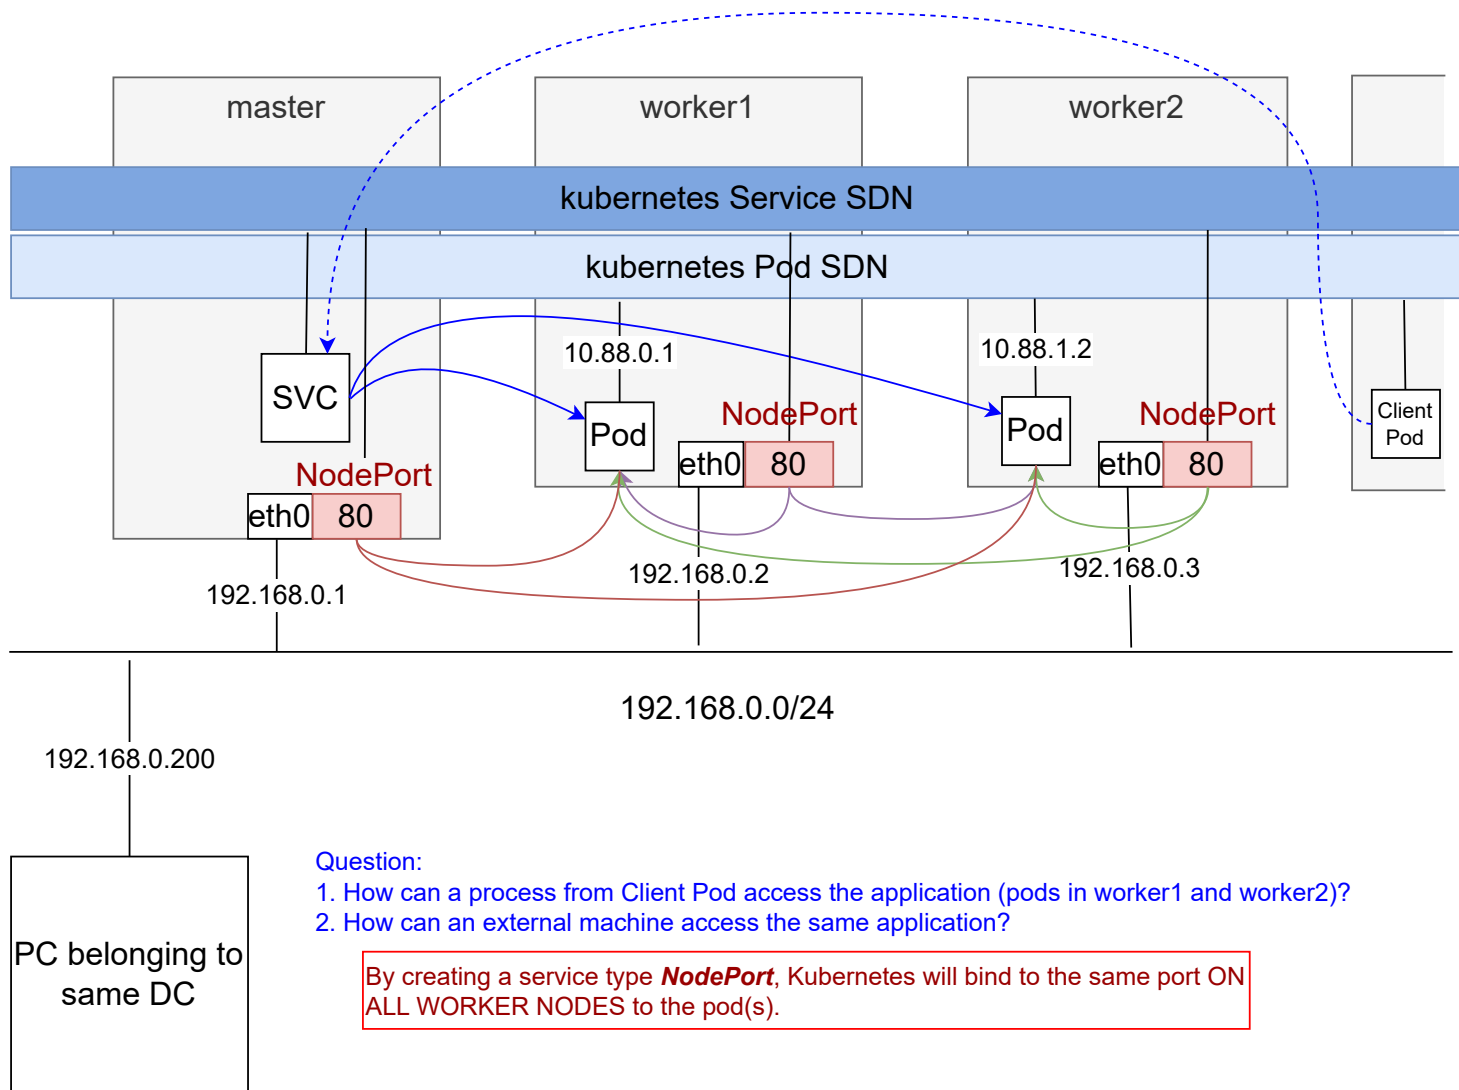
ROUTE

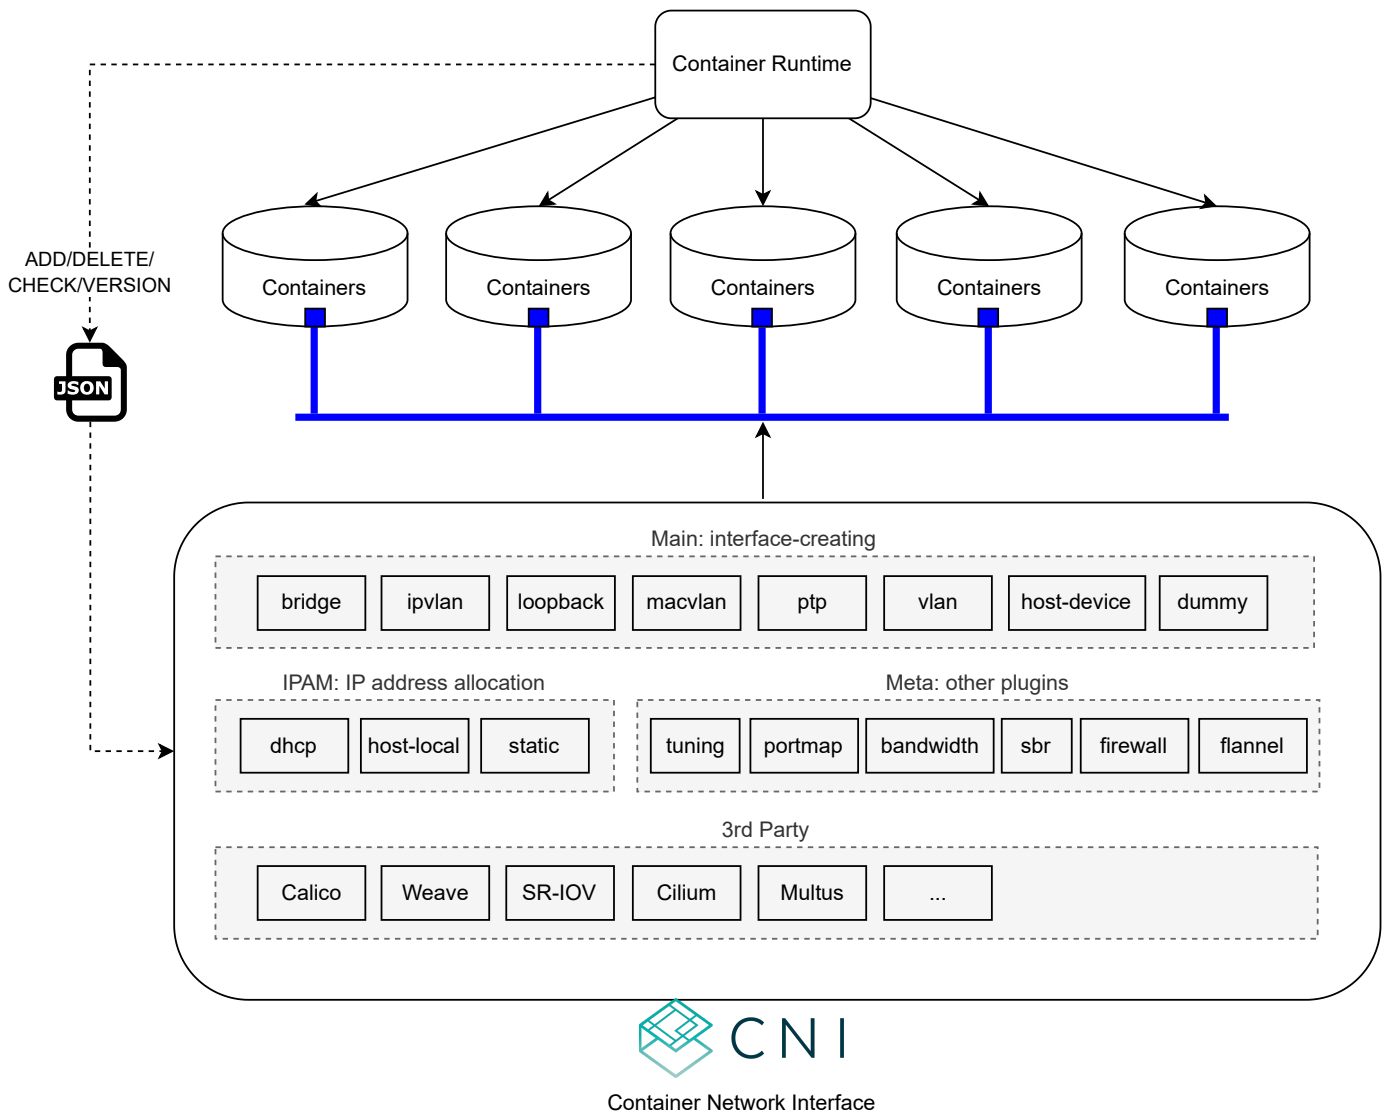
A route exposes the service to the external world.

Warning: A service "can" refer to different pods, if the pods have the same label.

Sample of how Services are used



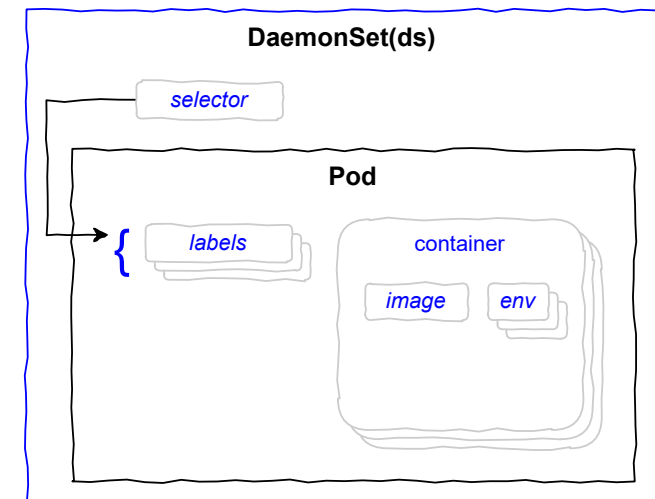
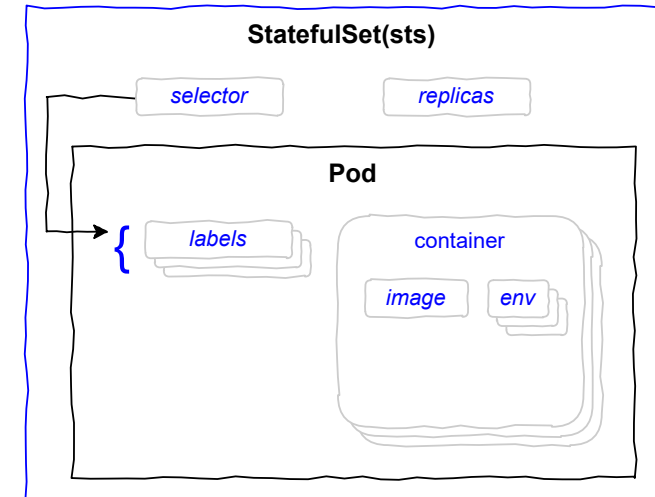




SERVICE(SVC)

```
oc expose <DC/DEPLOYMENT/RC/RS/POD> <RESOURCE_NAME>
```

```
DNS NAME = <SVC>.<PROJ>[.svc.cluster.local]  
ENVIRONMENT VARIABLE IN POD = <SVC>_SERVICE_HOST
```

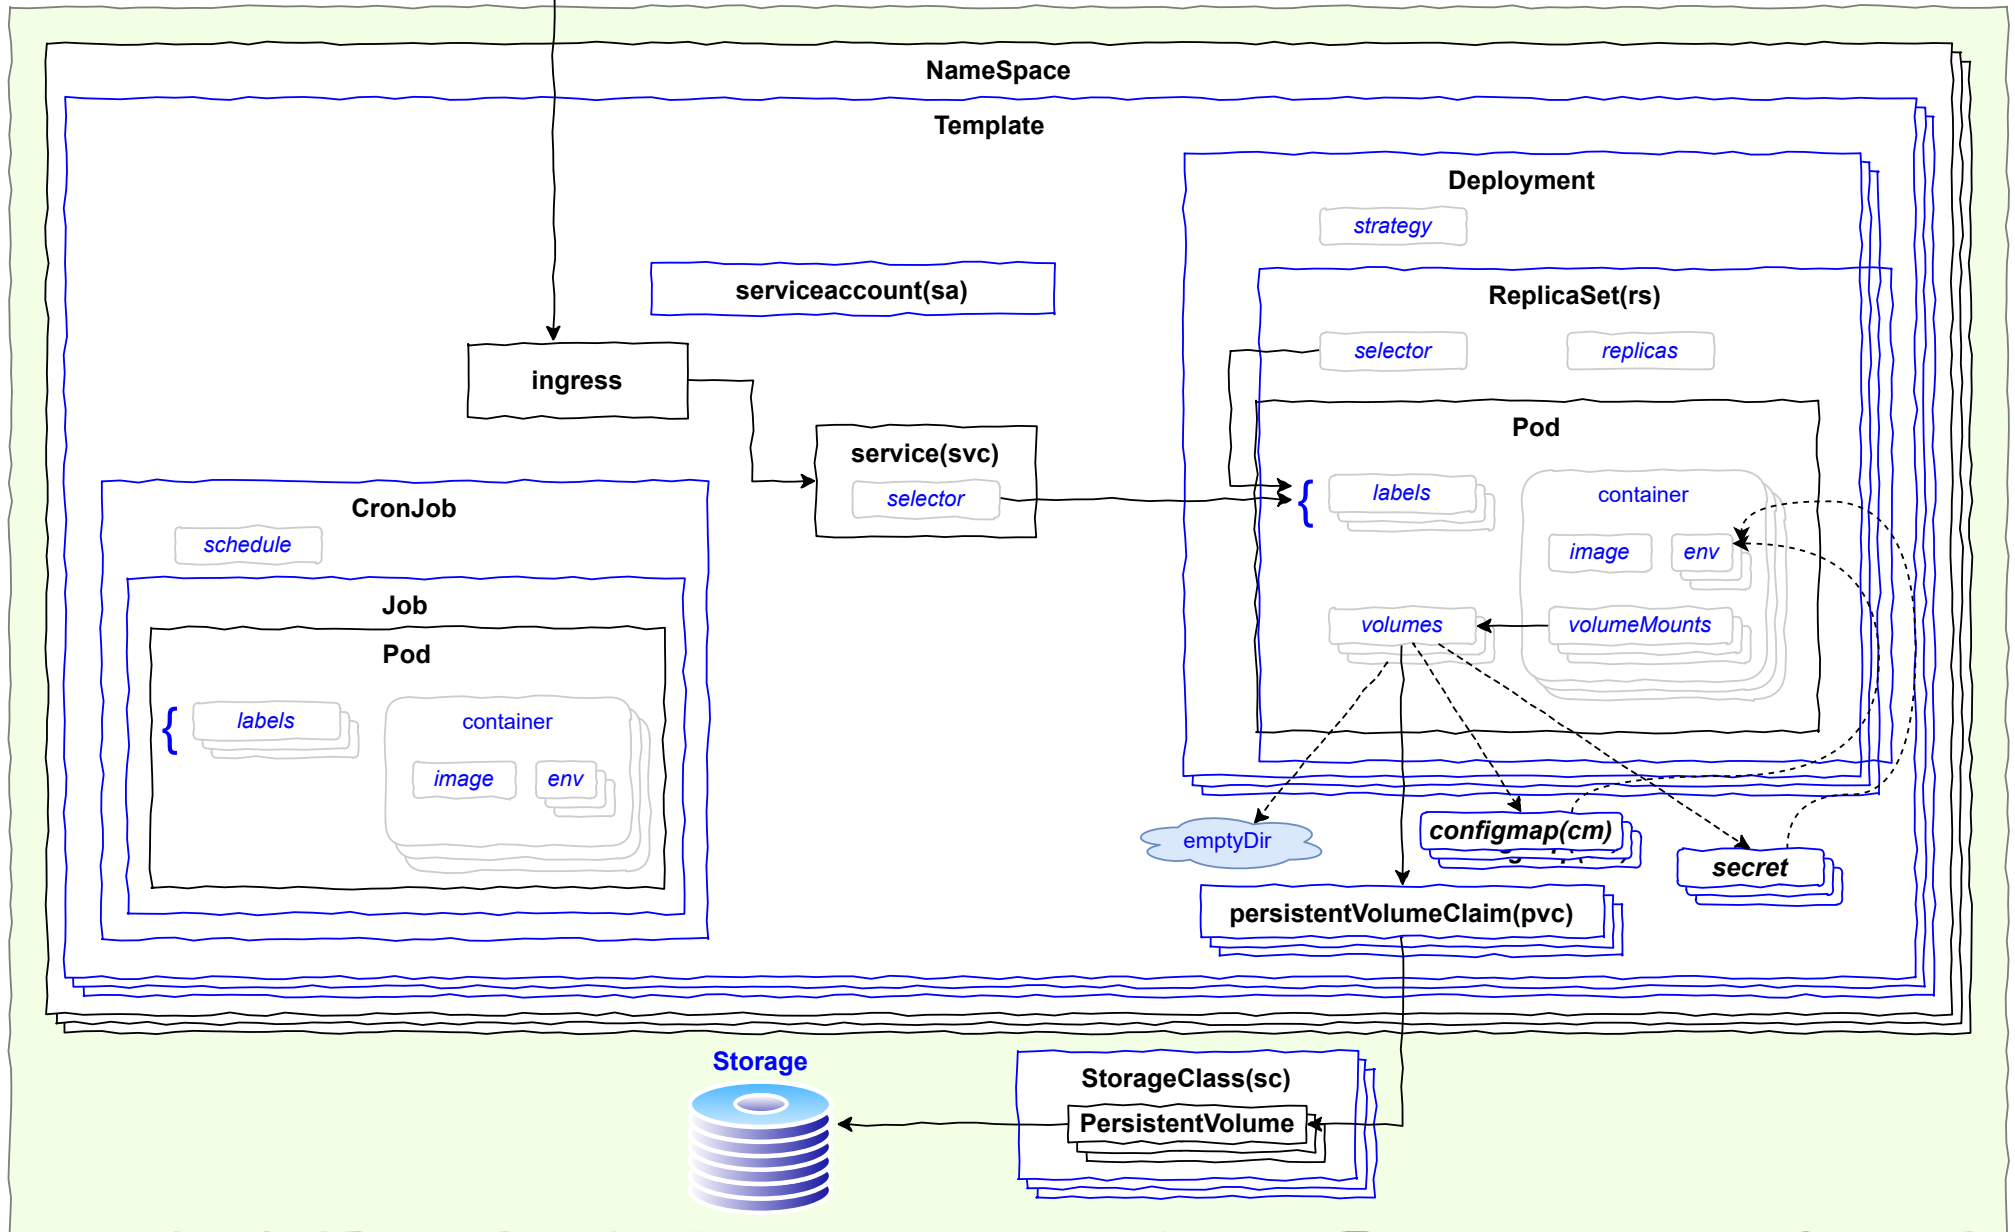


Kubernetes Resource Types

© 2025 Kelvin Lai



KUBERNETES CLUSTER



Template

apiVersion: template.openshift.io/v1

kind: Template

metadata:

name: mytemplate

annotations:

description: "Description"

objects:

- apiVersion: v1

kind: Pod

metadata:

name: \${APP_NAME}

spec:

containers:

- env:

- name: ACCESS_CODE

value: \${APP_PASS}

image: superapp/hyperimage

name: myApp

ports:

- containerPort: 8080

protocol: TCP

parameters:

- description: Name of Pod

name: POD_NAME

value: myPod

required: true

- description: Application Secret Access Code

name: APP_PASS

generate: expression

from: "[a-zA-Z0-9]{8}"

labels:

mylabel: myapp

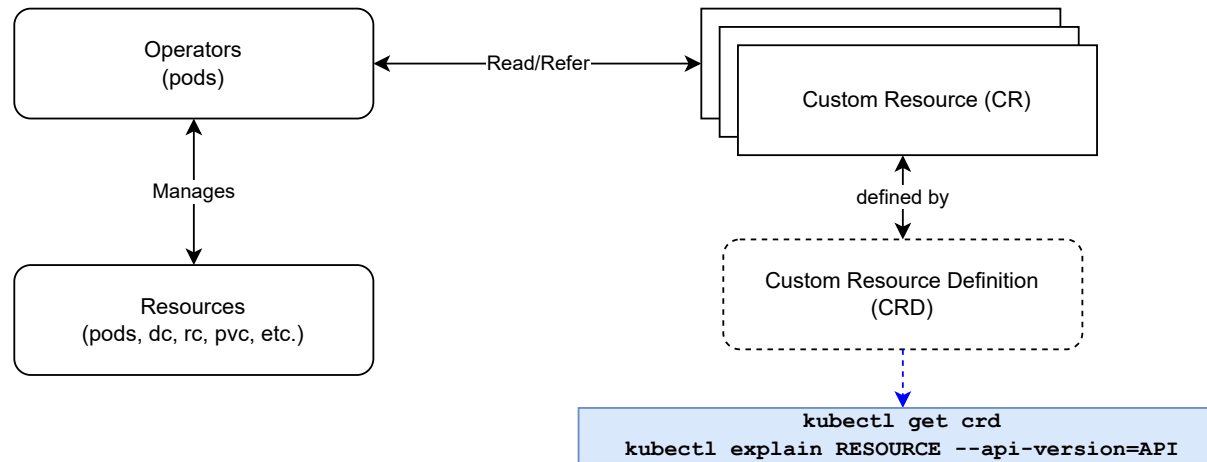
Object Creation Order
↓



```

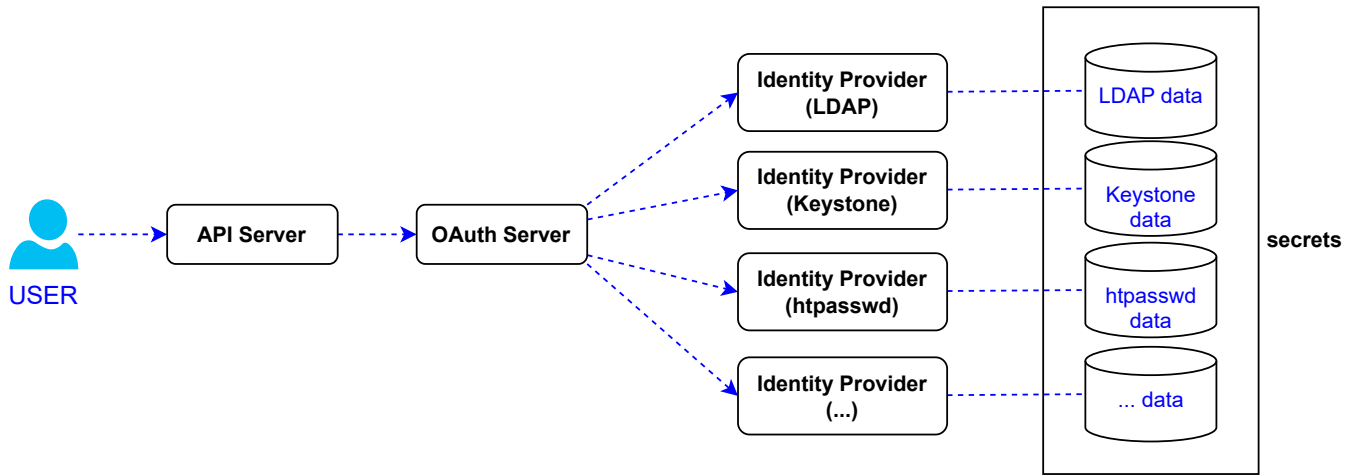
pipeline {
  options {                                // Global Options - can be overridden per stage basis
    timeout(time: 30, unit: 'MINUTES')
  }
  agent {                                  // Execution Context (where to run)
    node {
      label 'master'                       // IS used to run pipeline, (master - default minimal linux runtime)
    }
  }
  environment {                            // Global Vars
    DEV_PROJ = "super-app-v1"
    APP_NAME = "myapp"
  }
  stages {
    stage('stage 1') {
      steps {
        script {                           // DSL code (Groovy language code) must be embedded within script element
          openshift.withCluster() {
            openshift.withProject(env.DEV_PROJ) {
              openshift.selector("bc", "${APP_NAME}").startBuild("--wait=true", "--follow=true")
            }
          }
        }
      }
    }
    stage('stage 2') {                     // NOTE: project switches to default proj between stages.
      steps {
        sh 'echo hello from stage 2!'
      }
    }
    stage('manual approval') {
      steps {
        timeout(time: 60, unit: 'MINUTES') {
          input message: "Move to stage 3?"
        }
      }
    }
    stage('stage 3') {
      steps {
        sh 'echo hello from stage 3!. This is the last stage...'
      }
    }
  }
}

```

An operator is a control loop program and it can respond to events.
It communicates with the API server to manage k8s resources









```
$ oc get co | grep auth
$ oc get pods -n openshift-authentication-operator
```

Project: *openshift-authentication-operator*




Manages/Configures



Project: *openshift-authentication*

```
$ oc get pods -n openshift-authentication
```

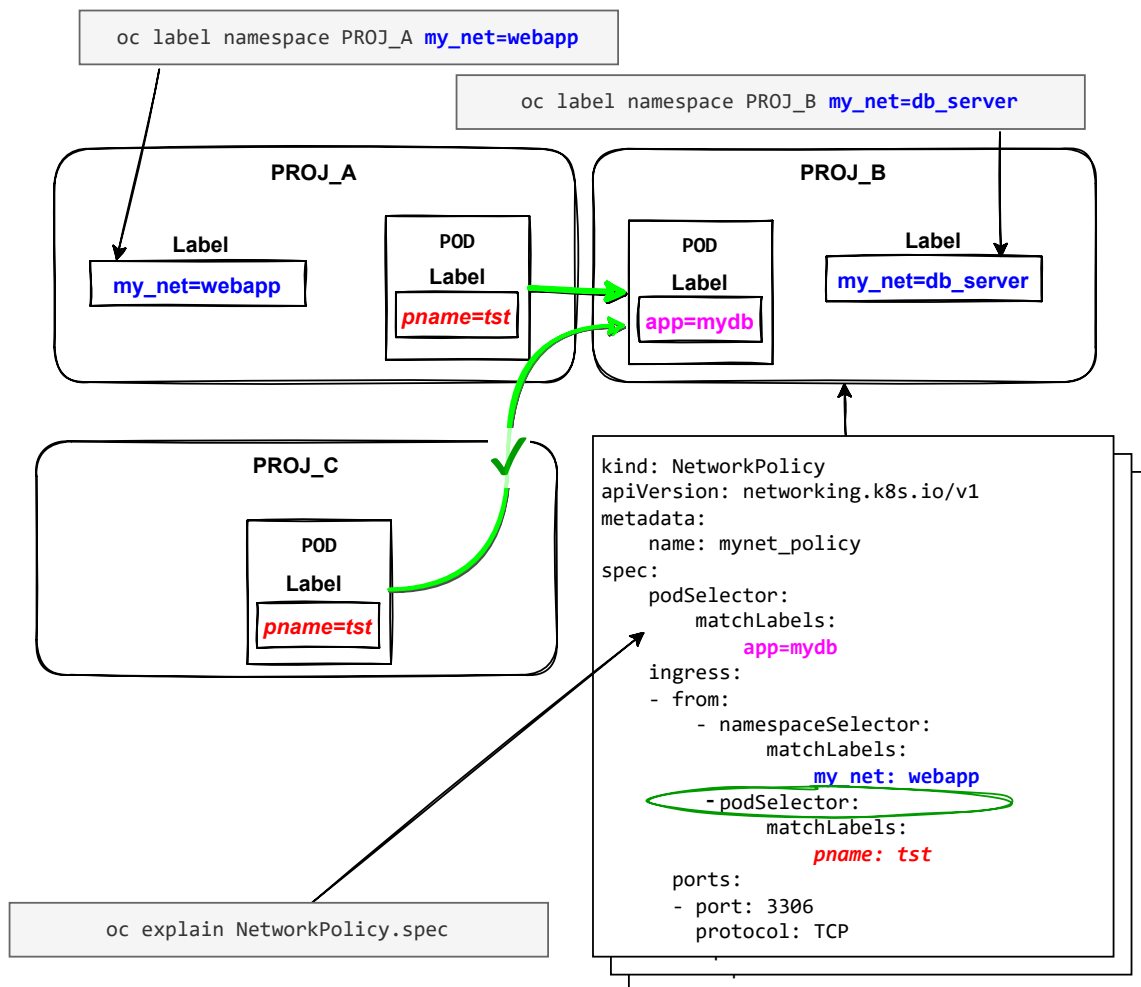

```
$ oc get oauth cluster -o yaml
apiVersion: config.openshift.io/v1
kind: OAuth
metadata:
  name: cluster
  ...
spec:
  identityProviders:
  - htpasswd:
      fileName:
        name: htpasswd-secret
      mappingMethod: claim
      name: htpasswd_provider
      type: HTTPasswd
```



Project: *openshift-config*


```
$ htpasswd -cBb ./myusers <USER> <PASSWORD>
$ oc create secret generic htpasswd-secret --from-file htpasswd=./myhtpasswd -n openshift-config
```

Termination Types



POD Scheduling

1. Get a list of all NODES
2. Go through all the predicates for **FILTERing**. If NODE fails predicate rule, remove from list. Region affinity.
3. With remainder list of NODES, **prioritize** them using the weightage rules. NO filtering of NODES done here. Zone anti-affinity.
4. Select the NODE with highest points.



```
kubectl label node <NODE> <KEY>=<VALUE>
```

Region

<KEY> = failure-domain.beta.kubernetes.io/region

A set of hosts in closed geographical area. High speed connectivity.

Zone (availability zone)

<KEY> = failure-domain.beta.kubernetes.io/zone

A set of hosts that share common critical infra components (ups, switch, storage)

Upgrade Path Graph: https://access.redhat.com/labs/ocpupgradegraph/update_channel

